

CSC 337



Homer
Simpson
Dan Castellan...



Bart Simpson
Nancy Cartwr...



Marge
Simpson
Julie Kavner



Lisa Simpson
Yeadley Smith

Multi-table SQL Queries (Joins)

Rick Mercer, Marty Stepp

Related tables and keys for Database simpsons

id	name	email
123	Bart	bart@fox.com
456	Milhouse	milhouse@fox.com
888	Lisa	lisa@fox.com
404	Ralph	ralph@fox.com

students

id	name
1234	Krabappel
5678	Hoover
9012	Obourn

teachers

id	name	teacher_id
10001	Computer Science 142	1234
10002	Computer Science 143	5678
10003	Computer Science 154	9012
10004	Informatics 100	1234

courses

student_id	course_id	grade
123	10001	B-
123	10002	C
456	10001	B+
888	10002	A+
888	10003	A+
404	10004	D+

grades

- **Primary key:** a column guaranteed to be unique for each record
 - e.g. Lisa's ID 888
- **Foreign key:** a column in table A storing a primary key value from table B
 - (e.g. records in **grades** with **student_id** of 888 are Lisa's grades)
- **Normalizing** splitting tables to improve structure / redundancy (linked by unique IDs)

SQL Joins

- A SQL join is a Structured Query Language (**SQL**) instruction to combine data from two or more sets of data (two or more tables)
- There are several types of joins
 - We'll use JOIN to select records that have matching values in both tables
 - Can match the primary key of one table with foreign key (s) in another table

General Form

SELECT column-names FROM table-name1

JOIN table-name2

ON column-name1 = column-name2

WHERE condition

- **JOIN** combines records from 2 or more tables that satisfy certain conditions
 - Creates a temporary table
- The **ON** clause specifies which records from each table are matched
- The rows are often linked by their key columns (id)
- The WHERE clause is optional

Easier to qualify table names

SELECT table1.col2, table2.col2, table1.col3 FROM table1name
JOIN table2name
ON table1.col1 = table2.col1
WHERE condition

- **Example**

```
SELECT students.name, grades.course_id, grades.grade FROM students  
JOIN grades  
ON students.id = grades.student_id  
WHERE students.name = 'Bart';
```

name	course_id	grade
Bart	10001	B-
Bart	10002	C

student_id	course_id	grade
123	10001	B-
123	10002	C
404	10002	B
404	10004	D+
456	10001	B+
456	10002	D-
888	10002	A+
888	10003	A+

grades

id	name	email	password
123	Bart	bart@fox.com	bartman
456	Milhouse	milhouse@fox.com	fallout
888	Lisa	lisa@fox.com	vegan
404	Ralph	ralph@fox.com	catfood

students

Filtered JOIN with ON and WHERE

```
SELECT students.name, grades.course_id, grades.grade FROM students
JOIN grades
ON students.id = grades.student_id
WHERE students.name = 'Bart';
```

```
+-----+-----+-----+
| name | course_id | grade |
+-----+-----+-----+
| Bart |      10001 | B-    |
| Bart |      10002 | C      |
+-----+-----+-----+
```

id	name	email	password
123	Bart	bart@fox.com	bartman
456	Milhouse	milhouse@fox.com	fallout
888	Lisa	lisa@fox.com	vegan
404	Ralph	ralph@fox.com	catfood

students

student_id	course_id	grade
123	10001	B-
123	10002	C
404	10002	B
404	10004	D+
456	10001	B+
456	10002	D-
888	10002	A+
888	10003	A+

grades

- **FROM / JOIN** glue tables together, **WHERE** filters results
- What goes in the **ON** clause, and what goes in **WHERE**?
 - **ON** directly links columns of the joined tables
 - **WHERE** sets additional constraints such as particular values (123, 'Bart')

Join Data from Two Tables

```
select * from grades;
```

student_id	course_id	grade
123	10001	B-
123	10002	C
456	10001	B+
888	10002	A+
888	10003	A+
404	10004	D+
404	10002	B
456	10002	D-

```
SELECT column-names FROM table-name1  
JOIN table-name2  
ON column-name1 = column-name2  
WHERE condition
```

```
select * from students;
```

id	name	email	password
123	Bart	bart@fox.com	bartman
404	Ralph	ralph@fox.com	catfood
456	Milhouse	milhouse@fox.com	fallout
888	Lisa	lisa@fox.com	vegan

- Let's get the student's names and grades where the grades are \geq 'B'
 - How many will there be?
 - What is the 1st record stored in a new table?

Answer and result

```
SELECT students.name, grades.grade FROM students
JOIN grades
ON students.id = grades.student_id
WHERE grades.grade > 'B';
```

name	grade
Bart	B-
Bart	C
Ralph	D+
Milhouse	B+
Milhouse	D-

JOIN data in grades and courses

```
select * from grades;
```

student_id	course_id	grade
123	10001	B-
123	10002	C
456	10001	B+
888	10002	A+
888	10003	A+
404	10004	D+
404	10002	B
456	10002	D-

```
SELECT column-names FROM table-name1
JOIN table-name2
ON column-name1 = column-name2
WHERE condition
```

```
select * from courses;
```

id	name	teacher_id
10001	Computer Science 142	1234
10002	Computer Science 143	5678
10003	Computer Science 190M	9012
10004	Informatics 100	1234

- All grades for Computer Science 143
 - id is 10001
 - join data from tables `courses` and `grades`
 - show `grades.student_id` `courses.name` and `grades.grade`

Answer and result

```
SELECT courses.name, grades.grade, grades.student_id FROM courses
JOIN grades
ON courses.id = grades.course_id
WHERE courses.id = '10001';
```

name	grade	student_id
Computer Science 142	B-	123
Computer Science 142	B+	456

Can we switch the
courses and
grades tables?

WHERE is Optional

```
select * from grades;
```

student_id	course_id	grade
123	10001	B-
123	10002	C
456	10001	B+
888	10002	A+

- Get all grades for everyone

```
SELECT * FROM grades
```

```
JOIN students
```

```
ON students.id = grades.student_id;
```

```
select * from students;
```

id	name	email	password
123	Bart	bart@fox.com	bartman
404	Ralph	ralph@fox.com	catfood
456	Milhouse	milhouse@fox.com	fallout
888	Lisa	lisa@fox.com	vegan

id	name	email	password	student_id	course_id	grade
123	Bart	bart@fox.com	bartman	123	10001	B-
123	Bart	bart@fox.com	bartman	123	10002	C
404	Ralph	ralph@fox.com	catfood	404	10004	D+
404	Ralph	ralph@fox.com	catfood	404	10002	B
456	Milhouse	milhouse@fox.com	fallout	456	10001	B+
456	Milhouse	milhouse@fox.com	fallout	456	10002	D-
888	Lisa	lisa@fox.com	vegan	888	10002	A+
888	Lisa	lisa@fox.com	vegan	888	10003	A+

Querying multi-table databases

- When we have larger datasets spread across multiple tables, we need queries that can answer high-level questions such as:
 - What courses has Bart taken and gotten a B- or better?
 - What courses have been taken by both Bart and Lisa?
 - Who are all the teachers Bart has had?
 - How many total students has Ms. Krabappel taught, and what are their names?
- To do this, we'll have to **JOIN** data from several tables
 - Add more JOIN clauses

Multi-way join

```
SELECT students.name, courses.name, grades.grade FROM courses  
JOIN grades ON grades.course_id = courses.id  
JOIN students ON grades.student_id = students.id  
WHERE grades.grade <= 'B-';
```

- More than two tables can be joined
- Data from which tables are joined?
 - _____
- What does the above query represent?
 - _____

What MariaDB would look like

```
MariaDB [simpsons]> SELECT students.name, courses.name, grades.grade FROM courses
-> JOIN grades ON grades.course_id = courses.id
-> JOIN students ON grades.student_id = students.id
-> WHERE grades.grade <= 'B-';
```

name	name	grade
Bart	Computer Science 142	B-
Ralph	Computer Science 143	B
Milhouse	Computer Science 142	B+
Lisa	Computer Science 143	A+
Lisa	Computer Science 190M	A+

5 rows in set (0.00 sec)

Designing a query:

- Figure out the proper SQL queries in the following way:
 - Which columns do I want in the result set? (**SELECT**)
 - Which table(s) contain the critical data? (**FROM**)
 - How are tables connected (**JOIN**)
 - How are values filtered (**WHERE**)?
- Test on a real data base in MariaDB (**imdb_small**)
- Write the PHP code to run those same queries
 - Fix all of those syntax errors

Our Lab 10 database: imdb

id	first_name	last_name	gender
433259	William	Shatner	M
797926	Britney	Spears	F
831289	Sigourney	Weaver	F
...			

actors

id	name	year	rank
112290	Fight Club	1999	8.5
209658	Meet the Parents	2000	7
210511	Memento	2000	8.7
...			

movies

actor_id	movie_id	role
433259	313398	Capt. James T. Kirk
433259	407323	Sgt. T.J. Hooker
797926	342189	Herself
...		

roles

movie_id	genre
209658	Comedy
313398	Action
313398	Sci-Fi
...	

movies_genres

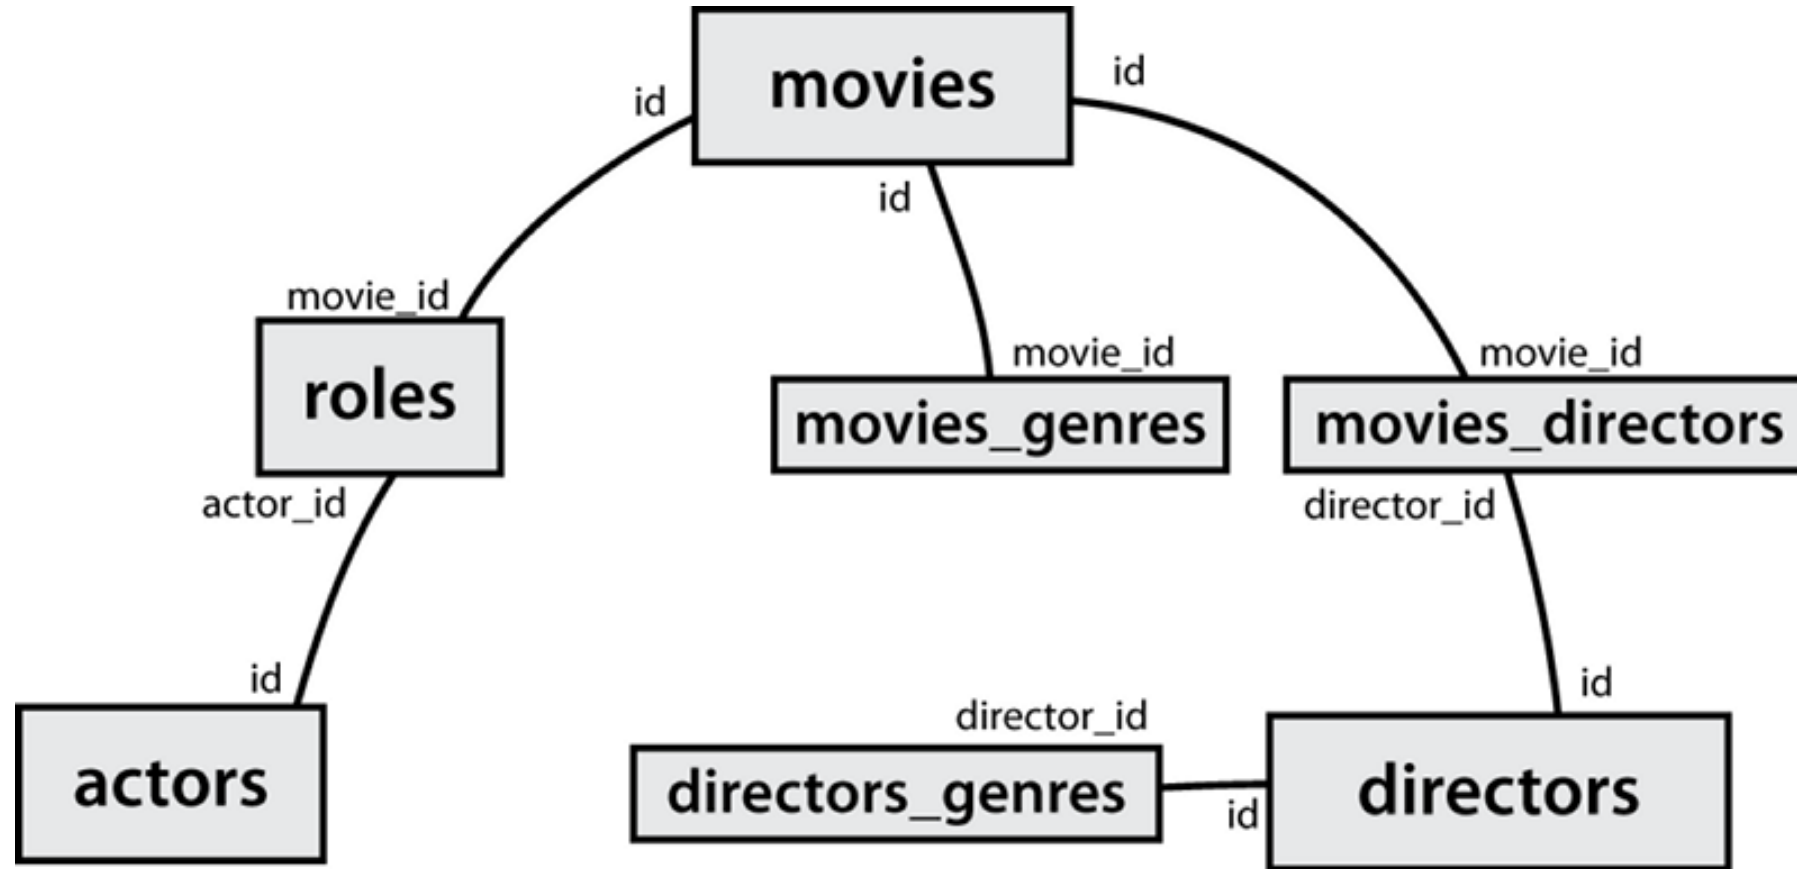
id	first_name	last_name
24758	David	Fincher
66965	Jay	Roach
72723	William	Shatner
...		

directors

director_id	movie_id
24758	112290
66965	209658
72723	313398
...	

movies_directors

IMDb table relationships / ids



Our other database: imdb_small

id	first_name	last_name	gender
433259	William	Shatner	M
797926	Britney	Spears	F
831289	Sigourney	Weaver	F
...			

actors

id	name	year	rank
112290	Fight Club	1999	8.5
209658	Meet the Parents	2000	7
210511	Memento	2000	8.7
...			

movies

actor_id	movie_id	role
433259	313398	Capt. James T. Kirk
433259	407323	Sgt. T.J. Hooker
797926	342189	Herself
...		

roles

movie_id	genre
209658	Comedy
313398	Action
313398	Sci-Fi
...	

movies_genres

id	first_name	last_name
24758	David	Fincher
66965	Jay	Roach
72723	William	Shatner
...		

directors

director_id	movie_id
24758	112290
66965	209658
72723	313398
...	

movies_directors

Possible queries

- How many people played a part in the movie "Lost in Translation"?
- What are the *names* of all the people who played a part in the movie "Lost in Translation"?
- What are the *names* of all movies Clint Eastwood has directed?
- What are the names of every actor who has appeared in a movie directed by Christopher Nolan?
- Practice with Marty Stepp's CSE 154 Query Tester
 - <http://www.martystepp.com/query/?username=cse154&password=cse154>
 - It can be a bit flaky